

Aravind Rajeswaran

Curriculum Vitae - September 2026
Website: <https://aravindr93.github.io>

Employment

Microsoft AI - Principal Researcher Grounding, SFT, and RL for computer use agents (CUA) Notable Projects: Fara1.5 , Fara-7B	<i>July 2025 - present</i>
Meta Platforms (FAIR) - Research Scientist Foundation models and benchmarks for Embodied AI Notable Projects: OpenEQA , VC-1 , R3M , RoboHive	<i>April 2021 - April 2025</i>
Google Brain - Research Internship Model-based reinforcement learning for robotics Notable Projects: MOREL , MBRL-Game	<i>June 2019 - May 2020</i>
OpenAI - Research Internship Algorithms for dexterous robot hands and multi-agent RL Notable Projects: DAPG	<i>June 2017 - Sep 2017</i>

Education

University of Washington Seattle Ph.D. in Computer Science & Engineering Advisers: Profs. Sham Kakade and Emo Todorov	<i>Sep 2016 – June 2021</i>
Indian Institute of Technology Madras , BTech (Hons.) Advisers: Profs. Balaraman Ravindran and Shankar Narasimhan	<i>Aug 2011 – July 2015</i>

Academic Awards

• Innovators Under 35 (TR35) from MIT Technology Review APAC.	<i>2026</i>
• Best paper award at the ICRA 2022 Scaling Robot Learning Workshop	<i>2022</i>
• Best paper award finalist at the RSS 2022 Scaling Robot Learning Workshop	<i>2022</i>
• J. P. Morgan PhD Fellowship in AI	<i>2020</i>
• Facebook PhD fellowship finalist in ML	<i>2020</i>
• Best paper award at IEEE SIMPAR	<i>2018</i>
• University of Washington PhD fellowship	<i>2016</i>
• Bhagyalakshmi and Krishna Ayengar award for best undergraduate thesis.	<i>2015</i>

Publications

Preprints

- [1] *Fara-1.5: Scalable Learning Environments for Computer Use Agents*
Note: alphabetical author ordering; Role: overall model training lead
A. Awadallah, S. Gupta, Y. Lara, Y. Lu, H. Mozannar, A. Nambi, Z. Nussbaum, Y. Pandya, A. Rajeswaran, C. Rosset, A. Taymanov, L. do Valle, V. Vineet, S. Whitehead, A. Zhao
arXiv Preprint, 2026.

Conference Articles

(* denotes equal first-author contribution, [†] denotes equal senior-authorship for leadership/advising)

- [2] *LOCATE 3D: Real-World Object Localization via Self-Supervised Learning in 3D*
P. McVay, S. Arnaud, A. Martin, A. Majumdar, K.M. Jatavallabhula, P. Thomas, R. Partsey, D. Dugas, A. Gejji, A. Sax, V.-P. Berges, M. Henaff, A. Jain, A. Cao, I. Prasad, M. Kalakrishnan, M. Rabbat, N. Ballas, M. Assran, O. Maksymets, A. Rajeswaran[†], F. Meier[†]
International Conference on Machine Learning (**ICML**), 2025. **(Spotlight)**
- [3] *OpenEQA: Embodied Question Answering in the Era of Foundation Models*
A. Majumdar, A. Ajay, X. Zhang, P. Putta, S. Yenamandra, M. Henaff, S. Silwal, P. McVay, O. Maksymets, S. Arnaud, K. Yadav, Q. Li, B. Newman, M. Sharma, V. Berges, S. Zhang, P. Agrawal, Y. Bisk, D. Batra, M. Kalakrishnan, F. Meier, C. Paxton, A. Sax, A. Rajeswaran
Computer Vision and Pattern Recognition (**CVPR**), 2024.
- [4] *MoDem-V2: Visuo-Motor World Models for Real-World Robot Manipulation*
P. Lancaster, N. Hansen, A. Rajeswaran, V. Kumar
International Conference on Robotics and Automation (**ICRA**), 2024.
- [5] *From LLMs to Actions: Latent Codes as Bridges in Hierarchical Robot Control*
Y. Shentu, P. Wu, A. Rajeswaran, P. Abbeel
International Conference on Intelligent Robots and Systems (**IROS**), 2024.
- [6] *RoboHive: A Unified Framework for Robot Learning*
V. Kumar, R. Shah, G. Zhou, V. Moens, V. Caggiano, A. Gupta, A. Rajeswaran
Advances in Neural Information Processing Systems (**NeurIPS**), 2023.
- [7] *Where are we in the search for an Artificial Visual Cortex for Embodied Intelligence?*
A. Majumdar, K. Yadav, S. Arnaud, J. Ma, C. Chen, S. Silwal, A. Jain, V.-P. Berges, T. Wu, J. Vakil, P. Abbeel, J. Malik, D. Batra, Y. Lin[†], O. Maksymets[†], A. Rajeswaran[†], F. Meier[†]
Advances in Neural Information Processing Systems (**NeurIPS**), 2023.
- [8] *Masked Trajectory Models for Prediction, Representation, and Control*
P. Wu, A. Majumdar, K. Stone, Y. Lin, I. Mordatch, P. Abbeel, A. Rajeswaran
International Conference on Machine Learning (**ICML**), 2023.
- [9] *On Pre-Training for Visuo-Motor Control: Revisiting a Learning-from-Scratch Baseline*
N. Hansen, Z. Yuan, Y. Ze, T. Mu, A. Rajeswaran[†], H. Su[†], H. Xu[†], X. Wang[†]
International Conference on Machine Learning (**ICML**), 2023.
- [10] *MoDem: Accelerating Visual Model-Based Reinforcement Learning with Demonstrations*
N. Hansen, Y. Lin, H. Su, X. Wang, V. Kumar, A. Rajeswaran
International Conference on Learning Representations (**ICLR**), 2023.
- [11] *Real World Offline Reinforcement Learning with Realistic Data Source*
G. Zhou, L. Ke, S. Srinivasa, A. Gupta, A. Rajeswaran, V. Kumar
International Conference on Robotics and Automation (**ICRA**), 2023.
- [12] *R3M: A Universal Visual Representation for Robot Manipulation*
S. Nair, A. Rajeswaran, V. Kumar, C. Finn, A. Gupta
Conference on Robot Learning (**CoRL**), 2022.
- [13] *The (Un)Surprising Effectiveness of Pre-Trained Vision Models for Control*
S. Parisi*, A. Rajeswaran*, S. Purushwalkam, A. Gupta
International Conference on Machine Learning (**ICML**), 2022. **(Long Oral)**
- [14] *CIC: Contrastive Intrinsic Control for Unsupervised Skill Discovery*
M. Laskin, H. Liu, X.B. Peng, D. Yarats, A. Rajeswaran, P. Abbeel
Advances in Neural Information Processing Systems (**NeurIPS**), 2022.

- [15] *Can Foundation Models Perform Zero-Shot Task Specification For Robot Manipulation?*
Y. Cui, S. Niekum, A. Gupta, V. Kumar, A. Rajeswaran
Learning for Dynamics and Control (**L4DC**), 2022.
- [16] *Decision Transformer: Reinforcement Learning via Sequence Modeling*
L. Chen, K. Lu, A. Rajeswaran, K. Lee, A. Grover, M. Laskin, P. Abbeel, A. Srinivas, I. Mordatch
Advances in Neural Information Processing Systems (**NeurIPS**), 2021.
- [17] *Visual Adversarial Imitation Learning using Variational Models*
R. Rafailov, T. Yu, A. Rajeswaran, C. Finn
Advances in Neural Information Processing Systems (**NeurIPS**), 2021.
- [18] *COMBO: Conservative Offline Model-Based Policy Optimization*
T. Yu, A. Kumar, R. Rafailov, A. Rajeswaran, S. Levine, C. Finn
Advances in Neural Information Processing Systems (**NeurIPS**), 2021.
- [19] *Reinforcement Learning with Latent Flow*
W. Shang, X. Wang, A. Srinivas, A. Rajeswaran, Y. Gao, P. Abbeel, M. Laskin
Advances in Neural Information Processing Systems (**NeurIPS**), 2021.
- [20] *Offline Reinforcement Learning from Images with Latent Space Models*
R. Rafailov, T. Yu, A. Rajeswaran, C. Finn
Learning for Dynamics and Control (**L4DC**), 2021.
- [21] *MOReL: Model-Based Offline Reinforcement Learning*
R. Kidambi*, A. Rajeswaran*, P. Netrapalli, T. Joachims
Advances in Neural Information Processing Systems (**NeurIPS**), 2020.
- [22] *A Game Theoretic Framework for Model Based Reinforcement Learning*
A. Rajeswaran, I. Mordatch, V. Kumar
International Conference on Machine Learning (**ICML**), 2020.
- [23] *Lyceum: An Efficient and Scalable Ecosystem for Robot Learning*
C. Summers, K. Lowrey, A. Rajeswaran, S. Srinivasa, E. Todorov
Learning for Dynamics and Control (**L4DC**), 2020.
- [24] *Meta-Learning with Implicit Gradients*
A. Rajeswaran*, C. Finn*, S. Kakade, S. Levine
Advances in Neural Information Processing Systems (**NeurIPS**), 2019.
- [25] *Online Meta-Learning*
C. Finn*, A. Rajeswaran*, S. Kakade, S. Levine
International Conference on Machine Learning (**ICML**), 2019.
- [26] *Plan Online, Learn Offline: Efficient Learning and Exploration via Model-Based Control*
K. Lowrey*, A. Rajeswaran*, S. Kakade, E. Todorov, I. Mordatch
International Conference on Learning Representations (**ICLR**), 2019.
- [27] *Dexterous Manipulation with Deep Reinforcement Learning: Efficient, General, and Low-Cost*
H. Zhu, A. Gupta, A. Rajeswaran, S. Levine, V. Kumar
International Conference on Robotics and Automation (**ICRA**), 2019.
- [28] *Learning Deep Visuomotor Policies for Dexterous Hand Manipulation*
D. Jain, A. Li, S. Singhal, A. Rajeswaran, V. Kumar, E. Todorov
International Conference on Robotics and Automation (**ICRA**), 2019.
- [29] *Reinforcement Learning for Non-Prehensile Manipulation: Transfer from Simulation to Physical System.*
K. Lowrey, S. Kolev, J. Dao, A. Rajeswaran, E. Todorov
IEEE International Conference on Simulation, Modeling, and Programming for Autonomous Robots (**SIMPAR**), 2018. **(Best Paper Award)**

- [30] *Variance Reduction for Policy Gradient with Action-Dependent Factorized Baselines*
C. Wu, A. Rajeswaran, Y. Duan, V. Kumar, A. Bayen, S. Kakade, I. Mordatch, P. Abbeel
International Conference on Learning Representations (**ICLR**), 2018. **(Full Oral)**
- [31] *Divide-and-Conquer Reinforcement Learning*
D. Ghosh, A. Singh, A. Rajeswaran, V. Kumar, S. Levine
International Conference on Learning Representations (**ICLR**), 2018.
- [32] *Learning Complex Dexterous Manipulation with Deep Reinforcement Learning and Demonstrations*
A. Rajeswaran*, V. Kumar*, A. Gupta, G. Vezzani, J. Schulman, E. Todorov, S. Levine
Proceedings of Robotics: Science and Systems (**RSS**), 2018.
- [33] *Towards Generalization and Simplicity in Continuous Control*
A. Rajeswaran*, K. Lowrey*, E. Todorov, S. Kakade
Advances in Neural Information Processing Systems (**NeurIPS**), 2017.
- [34] *EPOpt: Learning Robust Neural Network Policies Using Model Ensembles*
A. Rajeswaran, S. Ghotra, B. Ravindran, S. Levine
International Conference on Learning Representations (**ICLR**), 2017.

Journal Articles

- [35] *Identifying Topology of Power Distribution Networks Based on Smart Meter Data*
S. Jayadev, N. Bhatt, R. Pasumathy, A. Rajeswaran
IEEE Transactions on Smart Grid, 2017.
- [36] *A Graph Partitioning Approach for Leak Detection in Water Distribution Networks*
A. Rajeswaran, S. Narasimhan, S. Narasimhan
Computers & Chemical Engineering, 2017.

Media Coverage (Selected Subset)

- **VentureBeat** – “Microsoft’s Fara-7B is a computer-use AI agent that rivals GPT-4o and works directly on your PC” 2025
- **VentureBeat** – “Meta AI releases OpenEQA to spur ‘embodied intelligence’...” 2024
- **Yahoo Finance** – “Meta is teaching robots how to move on their own” 2023
- **SiliconANGLE** – “Meta’s AI researchers create artificial visual cortex for robots...” 2023
- **ZDNet** – “A Berkeley mash-up of AI approaches promises continuous learning” 2019

Mentoring

PhD Interns

- | | |
|---|--|
| <ul style="list-style-type: none"> • Philipp Wu (UC Berkeley, PhD → startup) • Arjun Majumdar (GATech, PhD → Meta) • Anurag Ajay (MIT, PhD → DeepMind) • Mandi Zhao (Stanford, PhD) • Nicklas Hansen (UCSD, PhD) | <ul style="list-style-type: none"> • Allan Zhou (Stanford, PhD → OpenAI) • Suraj Nair (Stanford, PhD → PI) • Liyiming Ke (UW, PhD → PI) • Andi Peng (MIT, PhD → startup) • Yuchen Cui (UT Austin, PhD → UCLA) |
|---|--|

University Students & AI Residents

- | | |
|--|---|
| <ul style="list-style-type: none"> • Shikhar Bahl (CMU, PhD → Skild.AI) • Aryan Jain (UC Berkeley, BS → Tesla) • Sergio Arnaud (AI Resident, FAIR → Meta) • Rafael Rafailov (Stanford, MS → Reflection AI) • Karmesh Yadav (FAIR → Georgia Tech, PhD) | <ul style="list-style-type: none"> • Ben Evans (UW, BS/MS → NYU, PhD) • Sarvjeet Ghotra (IITM Intern → MILA, PhD) • Colin Summers (UW, BS/MS → UW, PhD) • Catherine Cang (Berkeley, BS → Reflection AI) • Divye P. Jain (UW, BS/MS → Google) |
|--|---|

Professional Service and Teaching

Course Instructor and TA

- Fully designed and **taught** a special topics course at UW on deep RL for robotics. [[course website](#)]
- Teaching assistant for advanced graduate level machine learning courses at UW.

Workshops Organized

- Pretraining for Robot Learning ([website](#)), CoRL 2022.
- 3rd Offline RL workshop: Offline RL as a “Launchpad” ([website](#)), NeurIPS 2022.
- Object Representations for Learning and Reasoning ([website](#)), NeurIPS 2020.
- Generative Modeling and Model-Based Reasoning for Robotics and AI ([website](#)), ICML 2019.

Peer Review Service (Selected Subset)

- IEEE/CVF Conference on Computer Vision and Pattern Recognition – CVPR (2023)
- Neural Information Processing Systems – NeurIPS (2018, 2021, 2023, 2025, 2026)
- International Conference on Robotics and Automation – ICRA (2019, 2021)
- IEEE Robotics and Automation Letters (RA-L)